

# Information Processing 2 13

---

standard input/output, formatted input/output  
variadic functions, line input/output  
file access, error handling  
useful library functions

Ian Piumarta

Faculty of Engineering, KUAS

# last week: recursive structures, bit fields, unions

Chapter 6. Structures  
6.5 Self-referential Structures  
6.6 Table Lookup  
6.7 Typedef  
6.8 Unions  
6.9 Bit-fields

`sizeof()` and counting number of array elements  
using `typedef` for structures  
`union` and bit fields  
data structure: linked list, binary tree

# this week: input and output

`getchar()`, input redirection, pipes  
`putchar()`, output redirection, pipes  
`lower()`  
`printf()`  
variadic functions  
home-made `minprintf()`  
`scanf()`  
`FILE` access, `stdin`, `stdout`  
writing the `cat` program  
`stderr` and `exit()`  
line input and output  
string, char, memory functions

Chapter 7. Input and Output  
7.1 Standard input and output  
7.2 Formatted output – `printf`  
7.3 Variable-length argument lists  
7.4 Formatted input – `scanf`  
7.5 File access  
7.6 Error handling – `stderr` and `exit()`  
7.7 Line input and output  
7.8 Miscellaneous functions

## standard input and output

simplest input mechanism: `int getchar(void)`

- reads one character from the *standard input* (normally the keyboard)
- returns `EOF` when the end of file is reached

simplest output mechanism: `int putchar(int c)`

- writes one character to the *standard output* (normally a ‘terminal’ window)
- returns `c` or `EOF` if an error occurs

the source of input and destination of output can often be changed using *redirection*

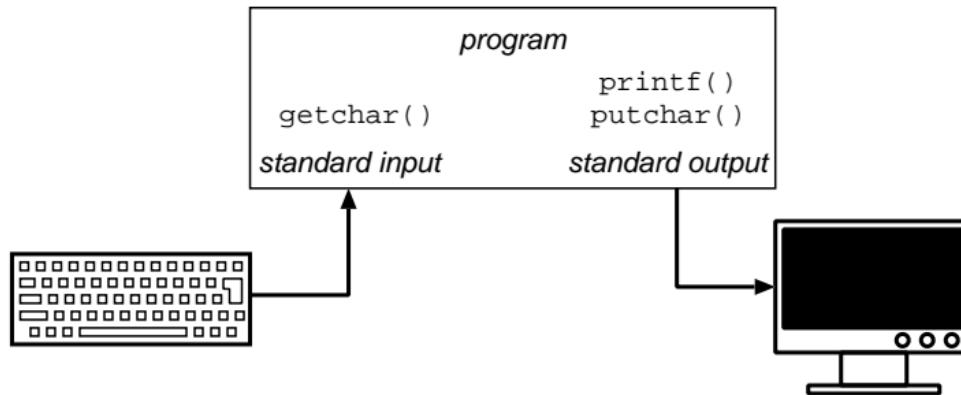
- on Linux (including WSL) and Mac this works as described below
- on Windows it depends on which shell or console you are using

## input from keyboard, output to screen

standard input is usually the keyboard

standard output is usually the screen (terminal/console window)

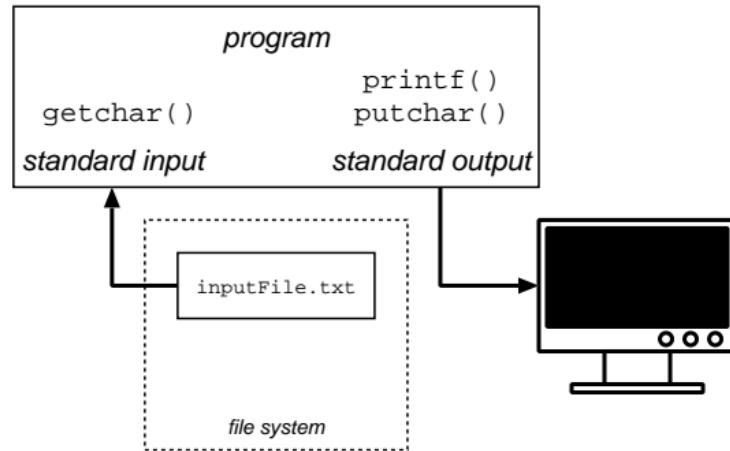
```
$ ./program
```



## input from a file, output to screen

standard input can be *redirected* to come from a file instead of from the keyboard

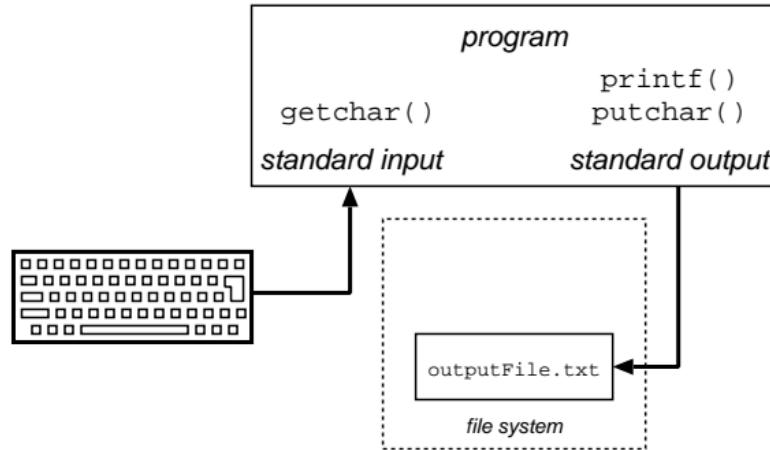
```
$ ./program < inputFile.txt
```



# input from keyboard, output to a file

standard output can be *redirected* to go to a file instead of to the screen

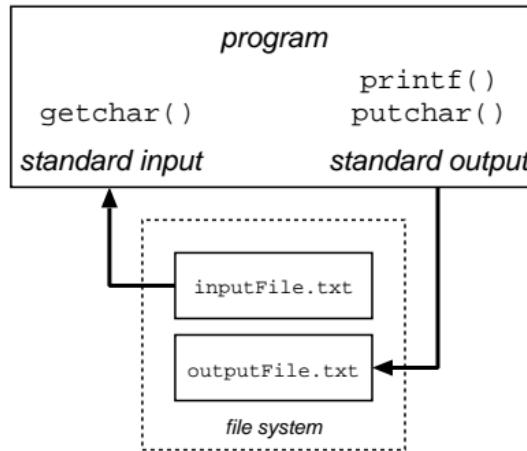
```
$ ./program > outputFile.txt
```



## input from a file, output to a file

standard input and output can be *redirected* at the same time

```
$ ./program < inputFile.txt > outputFile.txt
```



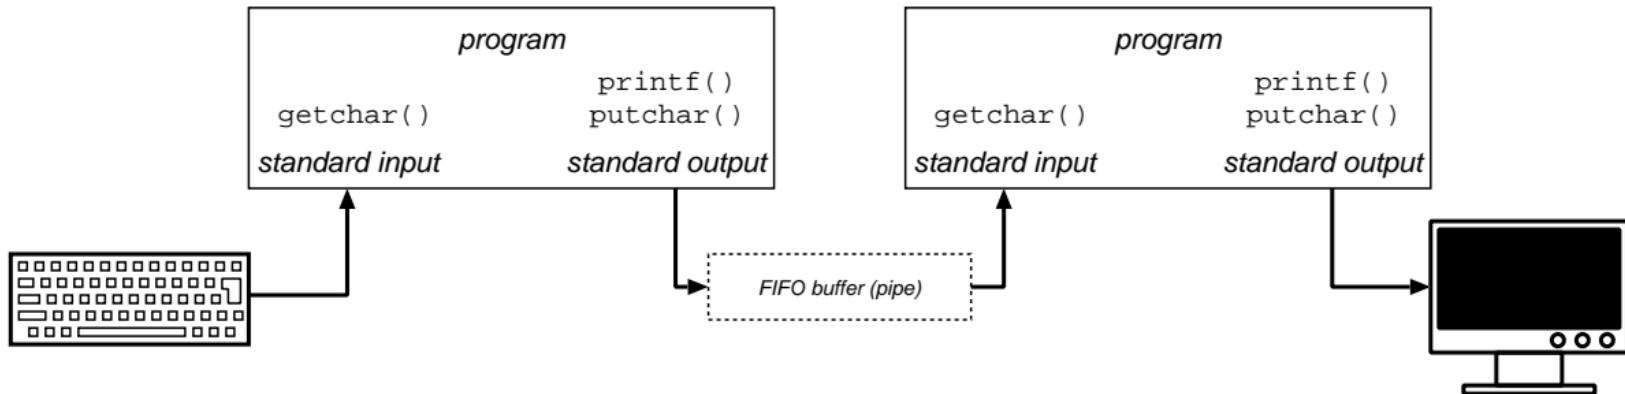
## piping output to input of another program

standard output can be *piped* into the standard input of another program

e.g: output generated by programA will be read as input by programB

a small *buffer* allows programA to continue writing even when programB is busy

```
$ ./programA | ./programB
```



## formatted output

general printing mechanism: `int printf(char *format, arguments...)`

- prints `format` on the *standard output*
- embeds zero or more *arguments*, according to '%' conversions in `format`

a conversion specifier contains: `% flags width . precision conversion`

*flags* are zero or more of

- left-justify the output (default is right-justified)
- + always print a sign character, either '+' or '-'
- (blank) leave a space before a positive number (instead of the '+' sign)
- 0 pad the output on the left with leading zeros (instead of spaces)

*width* is the minimum width of the output (padding will be added)

*precision* applied to string: maximum number of characters to print  
floating-point: number of digits after the decimal point  
integer: minimum number of digits

## formatted output

a conversion specifier contains:  $\% \text{ flags width . precision conversion}$

<i>conversion</i>	<i>type</i>	<i>representation</i>
d i	int	signed decimal integer
u	int	unsigned decimal number
o	int	unsigned octal integer (without leading zero)
x X	int	unsigned hexadecimal number (without leading 0x)
c	int	single character
s	char *	strings (nul-terminated character array)
f	double	$[-]m.ddd\ldots d$ where the number of <b>d</b> s is given by the precision
e E	double	$[-]m.ddd\ldots d \times 10^{\pm xx}$
g G	double	$\%e$ if the exponent is small, otherwise $\%f$
p	pointer	unsigned hexadecimal (usually)
%		print a literal '%' character

## formatted output

`int` size can be specified using a character before the conversion

- h short integer
- l long integer
- z the type appropriate for a `size_t` or `ssize_t` argument

*width* or *precision* can be specified as ‘\*’

- the value is read from the next argument, which must be an `int`

some useful applications of ‘\*’ using strings:

```
int n = 10;  
printf("%*s", n, ""); // print exactly n spaces  
printf("%.*s", n, s); // print at most n characters from string s
```

(the second example can be used to print `char` arrays that are not nul-terminated)

## formatted output examples

when printing the string “hello, world” (12 characters)

<i>format</i>	<i>output</i>
:%s:	:hello, world:
:%10s:	:hello, world:
:%.10s:	:hello, wor:
:% -10s:	:hello, world:
:%.15s:	:hello, world:
:%-15s:	:hello, world :
:%15s:	: hello, world:
:%15.10s:	: hello, wor:
:%-15.10s:	:hello, wor :

## formatted output hints and tips

beware of using `printf` to print a string without conversions

```
char *s = "unknown or unpredictable characters";  
printf(s);           // FAILS if s contains a % character  
printf("%s", s);   // correct and safe
```

`snprintf` is like `printf` except that it stores its output in a character array

```
int snprintf(char *array, size_t size, char *fmt, args...)
```

- `array` is where the formatted output string will be stored
- `size` is the size of the `array`
- `fmt` and `args...` are format string and arguments, as in `printf`

`snprintf` will not write more than `size` bytes into `array` (including nul terminator)

- return value is the number of bytes needed to store the entire output
- even if `size` was too small (which lets you grow `array` to the correct size)

## formatted output hints and tips

dynamically sizing the buffer for `snprintf`:

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    char s[] = "my string";
    int i = 42;

    char *buffer = 0; // buffer for snprintf output
    size_t size = 0; // size of buffer

    size = snprintf(buffer, size, "%i %s\n", i, s) + 1; // + 1 for NUL
    printf("allocate buffer of %zu bytes\n", size);
    buffer = malloc(size); // guaranteed to be correct size
    snprintf(buffer, size, "%i %s\n", i, s);
    printf("result: %s", buffer);
    free(buffer);

    return 0;
}
```

## variable-length arguments lists

`printf` accepts a variable number of arguments of unknown types

- but it can fetch those arguments correctly, under the control of the format string

functions that take a variable number of arguments are called *variadic*

the correct declaration for `printf` is: `int printf(char *format, ...)`

- the `int` result is the number of characters printed
- the `char *format` controls the number and types of the following arguments
- the `...` means the number and types of remaining arguments may vary  
i.e., the function is variadic

to recover the `...` arguments, `printf` uses several macros defined in `<stdarg.h>`

these macros allow `printf` to read each of the 'anonymous' arguments, one at a time

## variable-length arguments lists

### va\_list

- is a type, similar to a ‘pointer’ to the next argument
- used to declare a variable, conventionally called *ap*

### va\_start(*ap*, *argName*)

- initialises *ap* to point the argument following *argName*

### va\_arg(*ap*, *type*)

- fetches the next argument, which is assumed to be of the given *type*

### va\_end(*ap*)

- frees any resources being used by *ap*

```
void f(int a, int b, ...)  
    va_list ap;  
    va_start(ap, b);  
    while (more args) {  
        var = va_arg(ap, type);  
        ...  
    }  
    va_end(ap);  
}
```

## a minimal printf to demonstrate variadic functions

---

```
#include <stdio.h>
#include <stdarg.h>

int myprintf(char *fmt, ...)
{
    int n = 0;                                // number of characters printed
    va_list ap;                               // argument `pointer'
    va_start(ap, fmt);                      // point ap at the argument following fmt

    for (char *p = fmt;  *p;  ++p) {          // there is still more stuff to print
        if ('%' != *p || !p[1]) {            // not a conversion specification
            putchar(*p);                   // output the character verbatim
            continue;                      // finished with this character
        }
        switch (*++p) {
            case 'd': {
                int i = va_arg(ap, int);    // fetch next argument as an 'int'
                n += printf("%d", i);
                break;
            }
        }
    }
}
```

## a minimal printf to demonstrate variadic functions

---

```
case 'f': {
    double d = va_arg(ap, double); // fetch next argument as an 'double'
    n += printf("%f", d);
    break;
}
case 's': {
    char *s = va_arg(ap, char *);
    while (*s) ++n, putchar(*s++);
    break;
}
default:
    ++n, putchar(*p);
    break;
} // switch
} // for
va_end(ap); // release memory used by ap
return n;
}

int main()
{
    myprintf("%%<%s> [%d] {%f}\n",
             "hello", 42, 123.456);
    return 0;
}
```

## formatted input: scanf

the function `scanf` performs `printf`-like conversions in the opposite direction

```
int scanf(char *format, ...)
```

- reads characters from standard input, which must match `format`
- `format` can contain conversion specifications
- results of conversions are stored in variables using remaining *pointer* arguments
- conversion stops when the input does not match `format`
- return value is the number of successful conversions performed

```
#include <stdio.h>

int main()
{
    int i;
    while (1 == scanf("%d", &i)) // %d result needs a POINTER to an int variable
        printf("%d\n", i * i);
    return 0;
}
```

## formatted input: sscanf

<i>conversion</i>	<i>argument</i>	<i>effect</i>
d	int *	read a decimal integer
i	int *	read an integer in decimal, octal (leading 0), or hexadecimal (leading 0x)
o	int *	octal (without leading 0)
u	unsigned int *	unsigned decimal integer
x	int *	hexadecimal integer (without leading 0x)
c	char *	characters (without skipping whitespace)
s	char *	characters (skipping whitespace)
e f g	float *	floating-point number, option sign and exponent
%		matches a literal % character

integer conversions can be modified by h (short) or l (long)

floating-point conversions can be modified by l (double)

assignment can be suppressed with \* (pointer argument is omitted)

a field width for c and s limits the number of characters transferred

## scanning from a string

blanks in the format string skip any amount of input white space

- space, tab, newline, carriage return

the opposite of `snprintf`, the function `sscanf` reads input from a string

```
int sscanf(char *string, char *format, ...)
```

```
while (fgets(line, sizeof(line), stdin)) {
    int i[3];
    int n = sscanf(line, "%d %d %d", i, i+1, i+2);
    if (n < 1) break;
    switch (n) {
        case 1: printf("scalar %d\n", i[0]); break;
        case 2: printf("2d point %d, %d\n", i[0], i[1]); break;
        case 3: printf("3d point %d, %d, %d\n", i[0], i[1], i[2]); break;
    }
}
```

## file access

files other than standard input and output can be accessed and created

for example: `cat file1 file2 ...`

- open the next file named on the command line
- copy its contents to the standard output
- close the file

before accessing any file it must be *opened*

`fopen()` opens a file and returns a *file pointer* for subsequent file access

```
#include <stdio.h>

FILE *fp; // file pointer for access to the file
fp = fopen("test.txt", "r"); // can access test.txt contents using fp
```

note: **FILE** is a type name, not a structure tag

## fopen

fopen is declared as

```
FILE *fopen(char *path, char *mode)
```

- path is a filename (absolute, or relative to the working directory)
- mode is a string that specifies what kind of access is requested

"r"	open the file for reading only
"r+"	open the file for reading and writing
"w"	open or create the file for writing and truncate it
"w+"	open or create the file for reading and writing
"a"	open the file for writing and append to it

- the return value is either a pointer to a FILE (success) or 0 (error)

an open file is often called a *stream* (think: “flow of characters” in or out)

## basic input and output to files

`int getc(FILE *fp)` reads the next character from the stream `fp` and returns it, or `EOF` at end of file

`int putc(int c, FILE *fp)` writes the character `c` to the stream `fp` returning `c`, or `EOF` if an error occurs

when a C program is started, three files are automatically opened

`FILE *stdin` is connected to the standard input

`FILE *stdout` is connected to the standard output

`FILE *stderr` is connected to the standard error stream

(these are all declared in `<stdio.h>`)

`getchar` and `putchar` can be macros, defined using these functions and streams:

`#define getchar() getc(stdin)`

`#define putchar(c) putc(c, stdout)`

## formatted I/O using file streams

`printf` and `scanf` use `stdout` and `stdin` by default

several variants of them use an explicit file stream instead:

`int fscanf(FILE *fp, char *fmt, ...)` reads from `fp` instead of `stdin`

`int fprintf(FILE *fp, char *fmt, ...)` writes to `fp` instead of `stdout`

the function `int fclose(FILE *fp)` is the opposite of `fopen`

- it closes the file `fp`, making sure all data has been written to the file

we now have everything needed to write the cat program:

- for each command-line argument `a`
  - open `a` for reading as a file pointer `fp`
  - while the next character `c` read from `fp` is not `EOF`
    - \* write `c` to `stdout`
  - close the file pointer `fp`

## copying files using file streams

```
#include <stdio.h>

void filecopy(FILE *input, FILE *output)
{
    int c;
    while (EOF != (c = getc(input)))
        putc(c, output);
}

int main(int argc, char *argv[])
{
    for (int i = 1; i < argc; ++i) {
        char *path = argv[i];                // the file to copy to the output
        FILE *fp = fopen(path, "r");         // open for reading only
        if (!fp) {                          // open failed
            printf("%s: cannot open: %s\n", argv[0], path);
            return 1; // failure
        }
        filecopy(fp, stdout);              // copy the file to standard output
        fclose(fp);
    }

    return 0; // success
}
```

## error handling

the stream `stderr` is reserved for error messages

- `stderr` is an output stream, similar to `stdout`
- by default it is connected to the screen
- it remains connected to the screen even if `stdout` is redirected

the function `void exit(int status)` can be used to terminate the program

- it behaves like '`return status;`' does inside the `main` function

using these two facilities we can improve the error handling of `cat`

```
FILE *fp = fopen(path, "r");
if (!fp) {
    fprintf(stderr, "%s: cannot open: %s\n", argv[0], path); // use stderr for errors
    exit(1); // failure
}
```

(this code will continue to work even if moved out of `main` to another function)

## error handling

`int ferror(FILE *fp)` returns non-zero if an error has occurred on `fp`

`int feof(FILE *fp)` returns non-zero if end of file has occurred on `fp`

- note that the program must already have tried to read beyond the end of file
- this function does not tell you whether the next `getc(fp)` will return `EOF`

`void clearerr(FILE *fp)` clears any error (or end of file) condition on `fp`

## line-oriented input and output

```
char *fgets(char *line, int size, FILE *fp)
char *fputs(char *line, FILE *fp)
ssize_t getline(char **lineptr, size_t *linesize, FILE *fp)

#include <stdio.h> // getline()
#include <stdlib.h> // free()

int main() {
    char *line = 0;      // allocated and set by getline()
    size_t linesize = 0; // set by getline()
    while (getline(&line, &linesize, stdin) >= 0) // sets line and linesize
        printf("%zd: %s", linesize, line);
    if (line) free(line); // getline() allocates line by calling malloc()
    return 0;
}
```

## miscellaneous functions: ungetc

`int ungetc(int c, FILE *fp)` pushes the character `c` back onto the stream `fp`

this is similar to the `ungetchar(int c)` function we wrote

at least one character can be pushed back onto a stream

some systems allow more than one character to be pushed back

- but portable code should assume one character is the limit

`ungetc` works with any functions that read from a `FILE *`

- `scanf()`, `getc()`, `getchar()`, etc.

this allows your program to read ‘one character too far’ (e.g., to find the end of a line, a delimiter, etc.) and then ‘undo’ the reading of that one character

## miscellaneous functions: string operations

```
#include <string.h>

char *s, *t;
int c, n;

strcat(s, t)      concatenate t to end of s
strncat(s, t, n) concatenate at most n characters of t to end of s
strcmp(s, t)       return negative, zero, or positive for
                   s < t, s == t, or s > t, respectively
strncmp(s, t, n) same as strcmp() but only compares first n characters
strcpy(s, t)       copy t to s
strncpy(s, t, n)  copy at most n characters of t to s
strlen(s)         return length of s
strchr(s, c)      return pointer to first c in s, or NULL if not present
strrchr(s, c)     return pointer to last cx in s, or NULL if not present
```

## miscellaneous functions: character tests

```
#include <ctype.h>  
int c;
```

isalpha(c)	non-zero if <code>c</code> is alphabetic, 0 if not
isupper(c)	non-zero if <code>c</code> is upper case, 0 if not
islower(c)	non-zero if <code>c</code> is lower case, 0 if not
isdigit(c)	non-zero if <code>c</code> is digit, 0 if not
isalnum(c)	non-zero if <code>isalpha(c)</code> or <code>isdigit (c)</code> , 0 if not
isspace(c)	non-zero if <code>c</code> is blank, tab, newline, return, form-feed, vertical tab
toupper(c)	return <code>c</code> converted to upper case
tolower(c)	return <code>c</code> converted to lower case

## miscellaneous functions: mathematical operations

```
#include <math.h>
```

all angles are measured in radians ( $2\pi$  radians in a circle)

all arguments and results are `double`

<code>sin(x)</code>	sine of $x$
<code>cos (x)</code>	cosine of $x$
<code>atan2(y, x)</code>	arctangent of $\frac{y}{x}$
<code>exp(x)</code>	exponential function $e^x$
<code>log(x)</code>	natural (base $e$ ) logarithm of $x$ ( $x > 0$ )
<code>log10(x)</code>	common (base 10) logarithm of $x$ ( $x > 0$ )
<code>pow(x, y)</code>	$x^y$
<code>sqrt(x)</code>	square root of $x$ ( $x \geq 0$ )
<code>fabs(x)</code>	absolute value of $x$

## miscellaneous functions: storage management

```
#include <stdlib.h>
```

the type `void *` is used for ‘generic’ pointers (a pointer to any or unknown type)

- `void *` pointers behave as if they point to an object 1 byte large
- `void *` pointers can be converted to/from any other pointer type

`void *malloc(size_t n)` returns a pointer to a block of memory `n` bytes large

- the memory is *not* initialised (it contains random junk!)
- if the memory cannot be allocated, `0` (NULL) is returned

`void *calloc(size_t n, size_t size)` returns a pointer to memory for an array of `n` objects of the specified `size`

- the memory is cleared (filled with zeros), useful for allocating a `struct`  
`struct Foo *foop = calloc(1, sizeof(struct Foo));`

## miscellaneous functions: storage management

`void free(void *p)` frees the memory pointed to by `p`

- `p` must have been allocated by `malloc` or `calloc`
- memory can be freed in any order

`void *realloc(void *p, size_t size)` changes the size of the memory at `p`

- the memory at `p` is reallocated to contain `size` bytes of data
- the original contents of the memory are copied into the new memory
- `p` is `0` then `realloc` behaves like `malloc(size)`
- if new memory cannot be allocated then `0` (`NULL`) is returned
- if `p` is `0` (`NULL`) then `realloc` behaves like `malloc(size)`
- the memory should eventually be `free()`d, just as above

`realloc()` is very useful for buffers that have to grow larger on demand

## miscellaneous functions: storage management

it is an error to access memory after it has been freed

a typical error with freeing memory is:

```
for (p = head; p != 0; p = p->next) // WRONG
    free(p); // this invalidates the memory at *p
```

(*p* no longer points to valid memory when *p->next* is evaluated)

the correct way is to access *p->next* *before* freeing *p*:

```
for (p = head; p != 0; p = q) {
    q = p->next;
    free(p);
}
```

## miscellaneous functions: random number generation

```
#include <stdlib.h>
```

int **rand**(void) returns a random integer between 0 and **RAND\_MAX**

one way to produce random floating-point numbers in the interval [0, 1) is:

```
#define frand() ((double)rand() / ((double)RAND_MAX+1))
```

the sequence of random numbers is typically the same

**srand(unsigned)** can be used to ‘seed’ the generator for more random sequence

one way to do this (on Unix-like systems) is using the current time

```
#include <sys/time.h>
```

```
struct timeval tv;
```

```
gettimeofday(&tv, 0);
```

```
srand(tv.tv_usec); // microseconds part of current time
```

## miscellaneous functions: command execution

```
#include <stdlib.h>
```

int system(char \*s) executes the command in the string s

- this can be written exactly as you would enter it at a shell prompt
- the result is the exit status of the command s

```
system("date");
```

next week we look at more operating system services

# next week: operating system interface

`read()` and `write()`

`copy` program

`getchar()` function

`fcntl`

variadic `error()` function

implementation of `fopen()` and `getc()`

implementation of `fopen()`

`dirent` and directory listing

`fsize` and listing directories

diy storage allocator

## Chapter 8. The UNIX System Interface

8.1 File descriptors

8.2 Low-level I/O – `read()` and `write()`

8.3 `open()`, `creat()`, `close()`, `unlink()`

8.4 Random access – `lseek()`

8.5 Example – an implementation of `fopen()` and `getc()`

8.6 Example – listing directories

8.7 Example – a storage allocator

# assignment

please download assignment (projects) from  
MS Team “Information Processing 2”, “General” channel

## projects

complete any combination of projects, up to 50 points

[10] prime number generator

[10] Monte Carlo simulation to calculate  $\pi$  experimentally

[ 5] improve the RPN calculator program with `scanf`

[10] add variables to the RPN calculator program

[ 5] data type and arithmetic operations for  $3 \times 3$  matrices

[10] data type and arithmetic operations for 3-element vectors

[10] printing the contents of a file hierarchy

[ 5] searching for specific names in a file hierarchy

13:00 5  
13:05 5  
13:10 5  
13:15 5  
13:20 5  
13:25 5  
13:30 5  
13:35 5  
13:40 5  
13:45 5  
13:50 5  
13:55 5  
14:00 5  
14:05 5  
14:10 5  
14:15 5  
14:20 5  
14:25 5

14:30 10 *break*

14:40 90 exercises

16:10 00 end